

CS 188 Project 1

January 21, 2020

0.1 Introduction


Welcome to **CS188 - Data Science Fundamentals!** We plan on having you go through some grueling training so you can start crunching data out there... in today's day and age "data is the new oil" or perhaps "snake oil" nonetheless, there's a lot of it, each with different purity (so pure that perhaps you could feed off it for a life time) or dirty which then at that point you can either decide to dump it or try to weed out something useful (that's where they need you...)

In this project you will work through an example project end to end.

Here are the main steps:

1. Get the data
2. Visualize the data for insights
3. Preprocess the data for your machine learning algorithm
4. Select a model and train
5. Does it meet the requirements? Fine tune the model

Steps to Machine Learning



0.2 Working with Real Data

It is best to experiment with real-data as opposed to artificial datasets.

There are many different open datasets depending on the type of problems you might be interested in!

Here are a few data repositories you could check out: - [UCI Datasets](#) - [Kaggle Datasets](#) - [AWS Datasets](#)

Below we will run through an California Housing example collected from the 1990's.

0.3 Setup

```
[2]: import sys
assert sys.version_info >= (3, 5) # python>=3.5

import sklearn
assert sklearn.__version__ >= "0.20" # sklearn >= 0.20

import numpy as np #numerical package in python
import os
%matplotlib inline
import matplotlib.pyplot as plt #plotting package

# to make this notebook's output identical at every run
np.random.seed(42)

#matplotlib magic for inline figures
%matplotlib inline
import matplotlib # plotting library
import matplotlib.pyplot as plt

# Where to save the figures
ROOT_DIR = "."
IMAGES_PATH = os.path.join(ROOT_DIR, "images")
os.makedirs(IMAGES_PATH, exist_ok=True)

def save_fig(fig_name, tight_layout=True, fig_extension="png", resolution=300):
    '''
        plt.savefig wrapper. refer to
        https://matplotlib.org/3.1.1/api/_as_gen/matplotlib.pyplot.savefig.html
    '''
    path = os.path.join(IMAGES_PATH, fig_name + "." + fig_extension)
    print("Saving figure", fig_name)
    if tight_layout:
        plt.tight_layout()
    plt.savefig(path, format=fig_extension, dpi=resolution)

[3]: import os
import tarfile
import urllib
DATASET_PATH = os.path.join("datasets", "housing")
```

0.4 Intro to Data Exploration Using Pandas

In this section we will load the dataset, and visualize different features using different types of plots. Packages we will use: - **Pandas**: is a fast, flexible and expressive data structure widely used for

tabular and multidimensional datasets. - **Matplotlib**: is a 2d python plotting library which you can use to create quality figures (you can plot almost anything if you're willing to code it out!) - other plotting libraries: [seaborn](#), [ggplot2](#)

```
[4]: import pandas as pd
```

```
def load_housing_data(housing_path):  
    csv_path = os.path.join(housing_path, "housing.csv")  
    return pd.read_csv(csv_path)
```

```
[5]: housing = load_housing_data(DATASET_PATH) # we load the pandas dataframe  
housing.head() # show the first few elements of the dataframe  
              # typically this is the first thing you do  
              # to see how the dataframe looks like
```

```
[5]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	

	population	households	median_income	median_house_value	ocean_proximity
0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	496.0	177.0	7.2574	352100.0	NEAR BAY
3	558.0	219.0	5.6431	341300.0	NEAR BAY
4	565.0	259.0	3.8462	342200.0	NEAR BAY

A dataset may have different types of features - real valued - Discrete (integers) - categorical (strings)

The two categorical features are essentially the same as you can always map a categorical string/character to an integer.

In the dataset example, all our features are real valued floats, except ocean proximity which is categorical.

```
[6]: # to see a concise summary of data types, null values, and counts  
      # use the info() method on the dataframe  
housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 20640 entries, 0 to 20639  
Data columns (total 10 columns):  
longitude           20640 non-null float64  
latitude            20640 non-null float64  
housing_median_age  20640 non-null float64  
total_rooms         20640 non-null float64
```

```

total_bedrooms      20433 non-null float64
population          20640 non-null float64
households          20640 non-null float64
median_income       20640 non-null float64
median_house_value  20640 non-null float64
ocean_proximity     20640 non-null object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB

```

```

[7]: # you can access individual columns similarly
     # to accessing elements in a python dict
     housing["ocean_proximity"].head() # added head() to avoid printing many columns.
     ↪ .

```

```

[7]: 0    NEAR BAY
     1    NEAR BAY
     2    NEAR BAY
     3    NEAR BAY
     4    NEAR BAY
     Name: ocean_proximity, dtype: object

```

```

[8]: # to access a particular row we can use iloc
     housing.iloc[1]

```

```

[8]: longitude      -122.22
     latitude       37.86
     housing_median_age    21
     total_rooms        7099
     total_bedrooms      1106
     population         2401
     households         1138
     median_income       8.3014
     median_house_value  358500
     ocean_proximity     NEAR BAY
     Name: 1, dtype: object

```

```

[9]: # one other function that might be useful is
     # value_counts(), which counts the number of occurrences
     # for categorical features
     housing["ocean_proximity"].value_counts()

```

```

[9]: <1H OCEAN      9136
     INLAND       6551
     NEAR OCEAN   2658
     NEAR BAY     2290
     ISLAND        5
     Name: ocean_proximity, dtype: int64

```

```
[10]: # The describe function compiles your typical statistics for each
# column
housing.describe()
```

```
[10]:
```

	longitude	latitude	housing_median_age	total_rooms	\
count	20640.000000	20640.000000	20640.000000	20640.000000	
mean	-119.569704	35.631861	28.639486	2635.763081	
std	2.003532	2.135952	12.585558	2181.615252	
min	-124.350000	32.540000	1.000000	2.000000	
25%	-121.800000	33.930000	18.000000	1447.750000	
50%	-118.490000	34.260000	29.000000	2127.000000	
75%	-118.010000	37.710000	37.000000	3148.000000	
max	-114.310000	41.950000	52.000000	39320.000000	

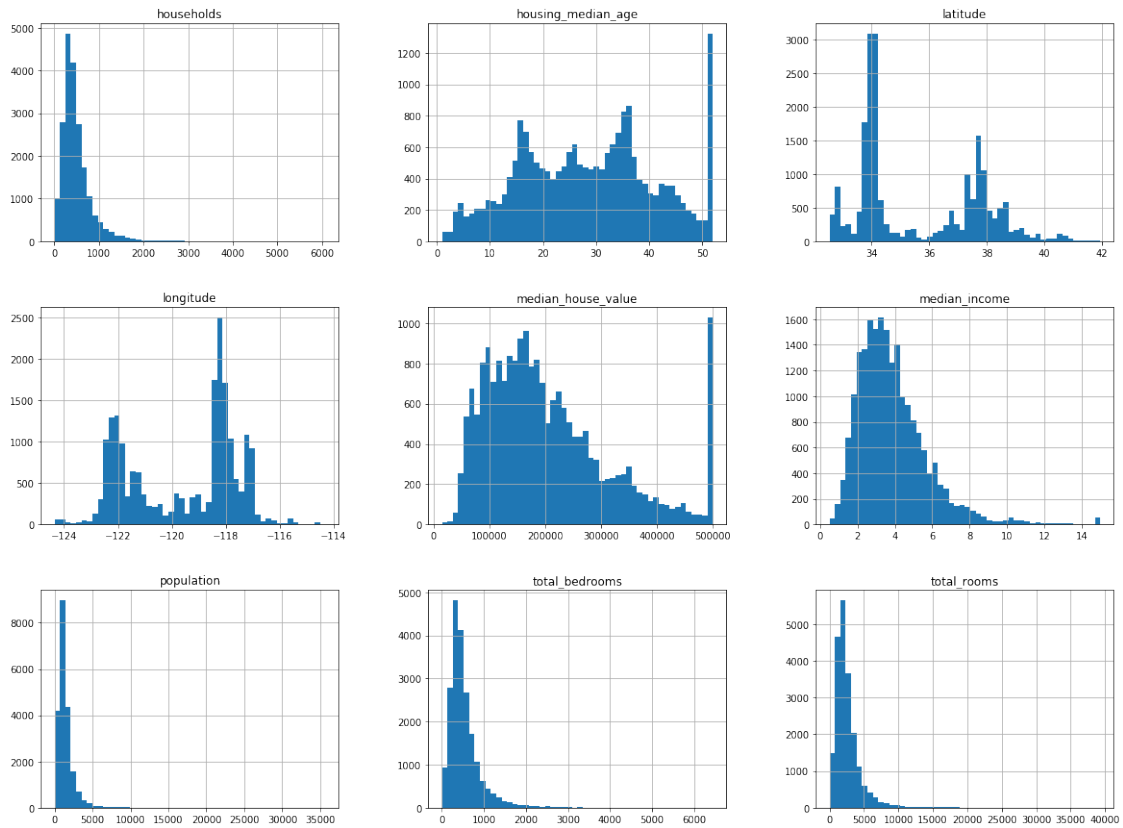
	total_bedrooms	population	households	median_income	\
count	20433.000000	20640.000000	20640.000000	20640.000000	
mean	537.870553	1425.476744	499.539680	3.870671	
std	421.385070	1132.462122	382.329753	1.899822	
min	1.000000	3.000000	1.000000	0.499900	
25%	296.000000	787.000000	280.000000	2.563400	
50%	435.000000	1166.000000	409.000000	3.534800	
75%	647.000000	1725.000000	605.000000	4.743250	
max	6445.000000	35682.000000	6082.000000	15.000100	

	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000


If you want to learn about different ways of accessing elements or other functions it's useful to check out the getting started section [here](#)

0.5 Let's start visualizing the dataset

```
[11]: # We can draw a histogram for each of the dataframes features
# using the hist function
housing.hist(bins=50, figsize=(20,15))
# save_fig("attribute_histogram_plots")
plt.show() # pandas internally uses matplotlib, and to display all the figures
# the show() function must be called
```



```
[12]: # if you want to have a histogram on an individual feature:
housing["median_income"].hist()
plt.show()
```



We can convert a floating point feature to a categorical feature by binning or by defining a set of intervals.

For example, to bin the households based on median_income we can use the `pd.cut` function


```
[13]: # assign each bin a categorical value [1, 2, 3, 4, 5] in this case.
housing["income_cat"] = pd.cut(housing["median_income"],
                               bins=[0., 1.5, 3.0, 4.5, 6., np.inf],
                               labels=[1, 2, 3, 4, 5])

housing["income_cat"].value_counts()
```

```
[13]: 3    7236
      2    6581
      4    3639
      5    2362
      1     822
      Name: income_cat, dtype: int64
```

```
[14]: housing["income_cat"].hist()
```


```
[14]: <matplotlib.axes._subplots.AxesSubplot at 0x1a24477d30>
```



Next let's visualize the household incomes based on latitude & longitude coordinates


```
[15]: ## here's a not so interesting way plotting it  
housing.plot(kind="scatter", x="longitude", y="latitude")  
save_fig("bad_visualization_plot")
```

Saving figure bad_visualization_plot



```
[16]: # we can make it look a bit nicer by using the alpha parameter,  
# it simply plots less dense areas lighter.  
housing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.1)  
save_fig("better_visualization_plot")
```

Saving figure better_visualization_plot



```
[17]: # A more interesting plot is to color code (heatmap) the dots
# based on income. The code below achieves this

# load an image of california
images_path = os.path.join('.', "images")
os.makedirs(images_path, exist_ok=True)
filename = "california.png"

import matplotlib.image as mpimg
california_img=mpimg.imread(os.path.join(images_path, filename))
ax = housing.plot(kind="scatter", x="longitude", y="latitude", figsize=(10,7),
                  s=housing['population']/100, label="Population",
                  c="median_house_value", cmap=plt.get_cmap("jet"),
                  colorbar=False, alpha=0.4,
                  )

# overlay the california map on the plotted scatter plot
# note: plt.imshow still refers to the most recent figure
# that hasn't been plotted yet.
plt.imshow(california_img, extent=[-124.55, -113.80, 32.45, 42.05], alpha=0.5,
           cmap=plt.get_cmap("jet"))
plt.ylabel("Latitude", fontsize=14)
plt.xlabel("Longitude", fontsize=14)
```

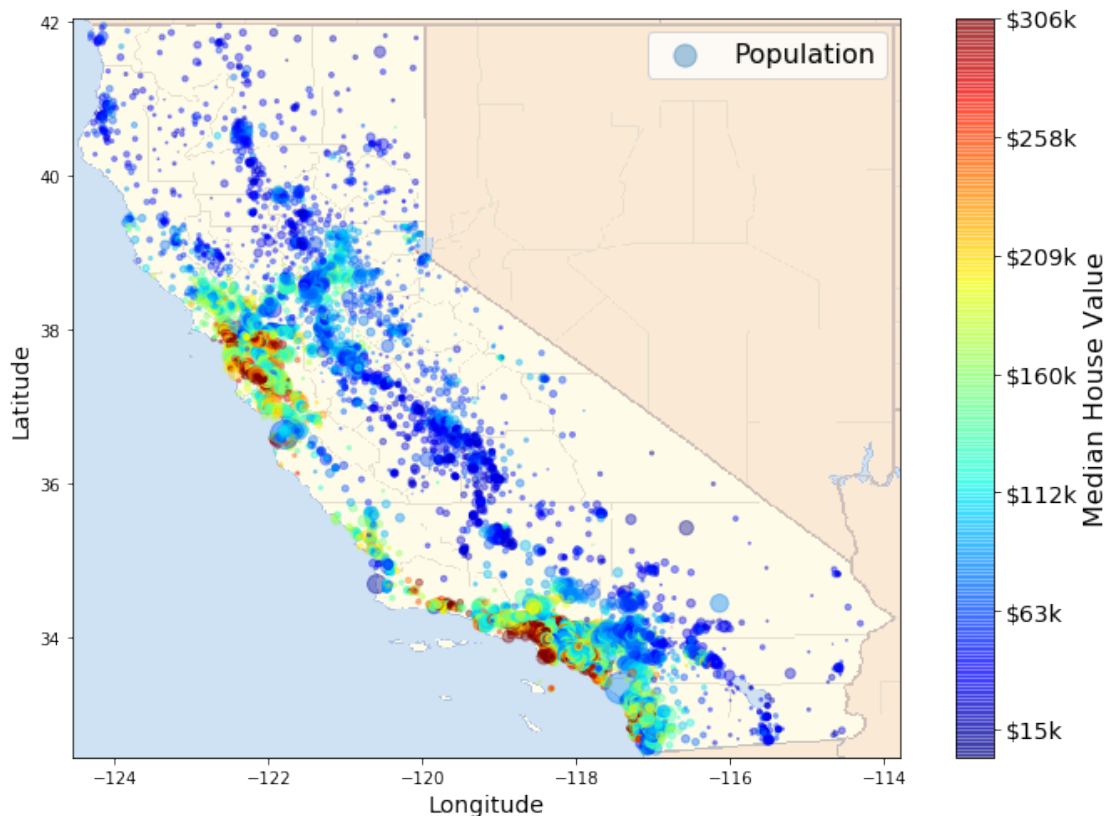
```

# setting up heatmap colors based on median_house_value feature
prices = housing["median_house_value"]
tick_values = np.linspace(prices.min(), prices.max(), 11)
cb = plt.colorbar()
cb.ax.set_yticklabels(["$%dk"%(round(v/1000)) for v in tick_values],
    ↪fontsize=14)
cb.set_label('Median House Value', fontsize=16)

plt.legend(fontsize=16)
save_fig("california_housing_prices_plot")
plt.show()

```

Saving figure california_housing_prices_plot



Not suprisingly, the most expensive houses are concentrated around the San Francisco/Los Angeles areas.

Up until now we have only visualized feature histograms and basic statistics.

When developing machine learning models the predictiveness of a feature for a particular target of intrest is what's important.

It may be that only a few features are useful for the target at hand, or features may need to be

augmented by applying certain transformations.

None the less we can explore this using correlation matrices.

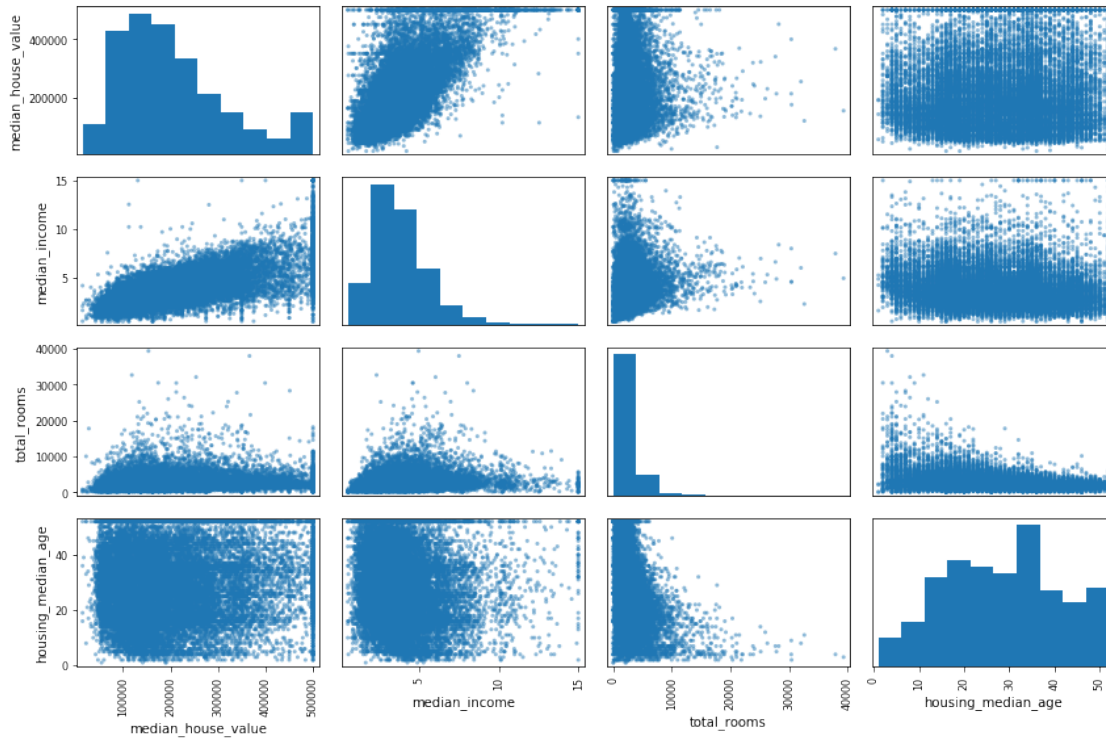
```
[18]: corr_matrix = housing.corr()
```

```
[19]: # for example if the target is "median_house_value", most correlated features ↵  
      ↪ can be sorted  
      # which happens to be "median_income". This also intuitively makes sense.  
      corr_matrix["median_house_value"].sort_values(ascending=False)
```

```
[19]: median_house_value    1.000000  
      median_income        0.688075  
      total_rooms          0.134153  
      housing_median_age    0.105623  
      households            0.065843  
      total_bedrooms        0.049686  
      population            -0.024650  
      longitude             -0.045967  
      latitude              -0.144160  
      Name: median_house_value, dtype: float64
```


```
[20]: # the correlation matrix for different attributes/features can also be plotted  
      # some features may show a positive correlation/negative correlation or  
      # it may turn out to be completely random!  
      from pandas.plotting import scatter_matrix  
      attributes = ["median_house_value", "median_income", "total_rooms",  
                   "housing_median_age"]  
      scatter_matrix(housing[attributes], figsize=(12, 8))  
      save_fig("scatter_matrix_plot")
```

Saving figure scatter_matrix_plot



```
[21]: # median income vs median house value plot plot 2 in the first row of top figure
housing.plot(kind="scatter", x="median_income", y="median_house_value",
              alpha=0.1)
plt.axis([0, 16, 0, 550000])
save_fig("income_vs_house_value_scatterplot")
```

Saving figure income_vs_house_value_scatterplot



0.5.1 Augmenting Features

New features can be created by combining different columns from our data set.

- $\text{rooms_per_household} = \text{total_rooms} / \text{households}$
- $\text{bedrooms_per_room} = \text{total_bedrooms} / \text{total_rooms}$
- etc.

```
[22]: housing["rooms_per_household"] = housing["total_rooms"]/housing["households"]
housing["bedrooms_per_room"] = housing["total_bedrooms"]/housing["total_rooms"]
housing["population_per_household"]=housing["population"]/housing["households"]
```

```
[23]: # obtain new correlations
corr_matrix = housing.corr()
corr_matrix["median_house_value"].sort_values(ascending=False)
```

```
[23]: median_house_value    1.000000
median_income             0.688075
rooms_per_household       0.151948
total_rooms               0.134153
housing_median_age        0.105623
households                0.065843
total_bedrooms            0.049686
```

```


population_per_household    -0.023737
population                  -0.024650
longitude                   -0.045967
latitude                   -0.144160
bedrooms_per_room          -0.255880
Name: median_house_value, dtype: float64

```

```

[24]: housing.plot(kind="scatter", x="rooms_per_household", y="median_house_value",
                    alpha=0.2)
plt.axis([0, 5, 0, 520000])
plt.show()

```



```

[25]: housing.describe()

```

```

[25]:
   count  longitude  latitude  housing_median_age  total_rooms  \
count    20640.000000    20640.000000    20640.000000    20640.000000
mean     -119.569704     35.631861        28.639486     2635.763081
std         2.003532         2.135952        12.585558     2181.615252
min      -124.350000     32.540000         1.000000         2.000000
25%      -121.800000     33.930000        18.000000     1447.750000
50%      -118.490000     34.260000        29.000000     2127.000000
75%      -118.010000     37.710000        37.000000     3148.000000
max       -114.310000     41.950000        52.000000    39320.000000

```

	total_bedrooms	population	households	median_income \
count	20433.000000	20640.000000	20640.000000	20640.000000
mean	537.870553	1425.476744	499.539680	3.870671
std	421.385070	1132.462122	382.329753	1.899822
min	1.000000	3.000000	1.000000	0.499900
25%	296.000000	787.000000	280.000000	2.563400
50%	435.000000	1166.000000	409.000000	3.534800
75%	647.000000	1725.000000	605.000000	4.743250
max	6445.000000	35682.000000	6082.000000	15.000100

	median_house_value	rooms_per_household	bedrooms_per_room \
count	20640.000000	20640.000000	20433.000000
mean	206855.816909	5.429000	0.213039
std	115395.615874	2.474173	0.057983
min	14999.000000	0.846154	0.100000
25%	119600.000000	4.440716	0.175427
50%	179700.000000	5.229129	0.203162
75%	264725.000000	6.052381	0.239821
max	500001.000000	141.909091	1.000000

	population_per_household
count	20640.000000
mean	3.070655
std	10.386050
min	0.692308
25%	2.429741
50%	2.818116
75%	3.282261
max	1243.333333

0.6 Preparing Dastaset for ML

Once we've visualized the data, and have a certain understanding of how the data looks like. It's time to clean!

Most of your time will be spent on this step, although the datasets used in this project are relatively nice and clean... it could get real dirty.

After having cleaned your dataset you're aiming for: - train set - test set

In some cases you might also have a validation set as well for tuning hyperparameters (don't worry if you're not familiar with this term yet..)

In supervised learning setting your train set and test set should contain (**feature**, **target**) tuples. - **feature**: is the input to your model - **target**: is the ground truth label - when target is categorical the task is a classification task - when target is floating point the task is a regression task

We will make use of [scikit-learn](#) python package for preprocessing.

Scikit learn is pretty well documented and if you get confused at any point simply look up the function/object!

```
[26]: from sklearn.model_selection import StratifiedShuffleSplit
# let's first start by creating our train and test sets
split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
for train_index, test_index in split.split(housing, housing["income_cat"]):
    train_set = housing.loc[train_index]
    test_set = housing.loc[test_index]

[27]: housing = train_set.drop("median_house_value", axis=1) # drop labels for
    ↪ training set features
                                           # the input to the model
    ↪ should not contain the true label
housing_labels = train_set["median_house_value"].copy()
```

0.6.1 Dealing With Incomplete Data

```
[28]: # have you noticed when looking at the dataframe summary certain rows
# contained null values? we can't just leave them as nulls and expect our
# model to handle them for us...
sample_incomplete_rows = housing[housing.isnull().any(axis=1)].head()
sample_incomplete_rows
```

```
[28]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
4629	-118.30	34.07	18.0	3759.0	NaN	
6068	-117.86	34.01	16.0	4632.0	NaN	
17923	-121.97	37.35	30.0	1955.0	NaN	
13656	-117.30	34.05	6.0	2155.0	NaN	
19252	-122.79	38.48	7.0	6837.0	NaN	

	population	households	median_income	ocean_proximity	income_cat	\
4629	3296.0	1462.0	2.2708	<1H OCEAN	2	
6068	3038.0	727.0	5.1762	<1H OCEAN	4	
17923	999.0	386.0	4.6328	<1H OCEAN	4	
13656	1039.0	391.0	1.6675	INLAND	2	
19252	3468.0	1405.0	3.1662	<1H OCEAN	3	

	rooms_per_household	bedrooms_per_room	population_per_household
4629	2.571135	NaN	2.254446
6068	6.371389	NaN	4.178817
17923	5.064767	NaN	2.588083
13656	5.511509	NaN	2.657289
19252	4.866192	NaN	2.468327

```
[29]: sample_incomplete_rows.dropna(subset=["total_bedrooms"])    # option 1: simply
      ↪ drop rows that have null values
```

```
[29]: Empty DataFrame
Columns: [longitude, latitude, housing_median_age, total_rooms, total_bedrooms,
population, households, median_income, ocean_proximity, income_cat,
rooms_per_household, bedrooms_per_room, population_per_household]
Index: []
```

```
[30]: sample_incomplete_rows.drop("total_bedrooms", axis=1)    # option 2: drop
      ↪ the complete feature
```

```
[30]:
```

	longitude	latitude	housing_median_age	total_rooms	population	\
4629	-118.30	34.07	18.0	3759.0	3296.0	
6068	-117.86	34.01	16.0	4632.0	3038.0	
17923	-121.97	37.35	30.0	1955.0	999.0	
13656	-117.30	34.05	6.0	2155.0	1039.0	
19252	-122.79	38.48	7.0	6837.0	3468.0	

	households	median_income	ocean_proximity	income_cat	\
4629	1462.0	2.2708	<1H OCEAN	2	
6068	727.0	5.1762	<1H OCEAN	4	
17923	386.0	4.6328	<1H OCEAN	4	
13656	391.0	1.6675	INLAND	2	
19252	1405.0	3.1662	<1H OCEAN	3	

	rooms_per_household	bedrooms_per_room	population_per_household
4629	2.571135	NaN	2.254446
6068	6.371389	NaN	4.178817
17923	5.064767	NaN	2.588083
13656	5.511509	NaN	2.657289
19252	4.866192	NaN	2.468327

```
[31]: median = housing["total_bedrooms"].median()
      sample_incomplete_rows["total_bedrooms"].fillna(median, inplace=True) # option
      ↪ 3: replace na values with median values
      sample_incomplete_rows
```

```
[31]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
4629	-118.30	34.07	18.0	3759.0	433.0	
6068	-117.86	34.01	16.0	4632.0	433.0	
17923	-121.97	37.35	30.0	1955.0	433.0	
13656	-117.30	34.05	6.0	2155.0	433.0	
19252	-122.79	38.48	7.0	6837.0	433.0	

	population	households	median_income	ocean_proximity	income_cat	\
4629	3296.0	1462.0	2.2708	<1H OCEAN	2	

6068	3038.0	727.0	5.1762	<1H OCEAN	4
17923	999.0	386.0	4.6328	<1H OCEAN	4
13656	1039.0	391.0	1.6675	INLAND	2
19252	3468.0	1405.0	3.1662	<1H OCEAN	3

	rooms_per_household	bedrooms_per_room	population_per_household
4629	2.571135	NaN	2.254446
6068	6.371389	NaN	4.178817
17923	5.064767	NaN	2.588083
13656	5.511509	NaN	2.657289
19252	4.866192	NaN	2.468327

Could you think of another plausible imputation for this dataset? (Not graded)

0.6.2 Prepare Data

```
[32]: # This cell implements the complete pipeline for preparing the data
# using sklearn's TransformerMixins
# Earlier we mentioned different types of features: categorical, and floats.
# In the case of floats we might want to convert them to categories.
# On the other hand categories in which are not already represented as integers
# → must be mapped to integers before
# feeding to the model.

# Additionally, categorical values could either be represented as one-hot
# → vectors or simple as normalized/unnormalized integers.
# Here we encode them using one hot vectors.

from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer

from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder

from sklearn.base import BaseEstimator, TransformerMixin

imputer = SimpleImputer(strategy="median") # use median imputation for missing
# → values
housing_num = housing.drop("ocean_proximity", axis=1) # remove the categorical
# → feature
# column index
rooms_ix, bedrooms_ix, population_ix, households_ix = 3, 4, 5, 6

#
```

```

class AugmentFeatures(BaseEstimator, TransformerMixin):
    """
    implements the previous features we had defined
    housing["rooms_per_household"] = housing["total_rooms"] /
    ↪ housing["households"]
    housing["bedrooms_per_room"] = housing["total_bedrooms"] /
    ↪ housing["total_rooms"]
    housing["population_per_household"] = housing["population"] /
    ↪ housing["households"]
    """
    def __init__(self, add_bedrooms_per_room = True):
        self.add_bedrooms_per_room = add_bedrooms_per_room
    def fit(self, X, y=None):
        return self # nothing else to do
    def transform(self, X):
        rooms_per_household = X[:, rooms_ix] / X[:, households_ix]
        population_per_household = X[:, population_ix] / X[:, households_ix]
        if self.add_bedrooms_per_room:
            bedrooms_per_room = X[:, bedrooms_ix] / X[:, rooms_ix]
            return np.c_[X, rooms_per_household, population_per_household,
                          bedrooms_per_room]
        else:
            return np.c_[X, rooms_per_household, population_per_household]

attr_adder = AugmentFeatures(add_bedrooms_per_room=False)
housing_extra_attribs = attr_adder.transform(housing.values)

num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('attribs_adder', AugmentFeatures()),
    ('std_scaler', StandardScaler()),
])

housing_num_tr = num_pipeline.fit_transform(housing_num)
numerical_features = list(housing_num)
categorical_features = ["ocean_proximity"]

full_pipeline = ColumnTransformer([
    ("num", num_pipeline, numerical_features),
    ("cat", OneHotEncoder(), categorical_features),
])

housing_prepared = full_pipeline.fit_transform(housing)

```

0.6.3 Select a model and train

Once we have prepared the dataset it's time to choose a model.

As our task is to predict the median_house_value (a floating value), regression is well suited for this.

```
[33]: from sklearn.linear_model import LinearRegression

lin_reg = LinearRegression()
lin_reg.fit(housing_prepared, housing_labels)

# let's try the full preprocessing pipeline on a few training instances
data = test_set.iloc[:5]
labels = housing_labels.iloc[:5]
data_prepared = full_pipeline.transform(data)

print("Predictions:", lin_reg.predict(data_prepared))
print("Actual labels:", list(labels))
```

```
Predictions: [425717.48517515 267643.98033218 227366.19892733 199614.48287493
161425.25185885]
```

```
Actual labels: [286600.0, 340600.0, 196900.0, 46300.0, 254500.0]
```

We can evaluate our model using certain metrics, a fitting metric for regression is the mean-squared-loss

$$L(\hat{Y}, Y) = \sum_i^N (\hat{y}_i - y_i)^2$$

where \hat{y} is the predicted value, and y is the ground truth label.

```
[34]: from sklearn.metrics import mean_squared_error

preds = lin_reg.predict(housing_prepared)
mse = mean_squared_error(housing_labels, preds)
rmse = np.sqrt(mse)
rmse
```

```
[34]: 67784.32202861732
```

1 TODO: Applying the end-end ML steps to a different dataset.

We will apply what we've learnt to another dataset (airbnb dataset). We will predict airbnb price based on other features.

2 [25 pts] Visualizing Data

2.0.1 [5 pts] Load the data + statistics

- load the dataset
- display the first few rows of the data
- drop the following columns: name, host_id, host_name, last_review
- display a summary of the statistics of the loaded data
- plot histograms for 3 features of your choice

```
[35]: DATASET_PATH = os.path.join("datasets", "airbnb")

def load_airbnb_data(airbnb_path):
    csv_path = os.path.join(airbnb_path, "AB_NYC_2019.csv")
    return pd.read_csv(csv_path)
airbnb = load_airbnb_data(DATASET_PATH)
airbnb = airbnb.drop(columns=['name', 'host_id', 'host_name', 'last_review'])
airbnb.head()
```

```
[35]:      id neighbourhood_group neighbourhood  latitude  longitude \
0   2539             Brooklyn   Kensington  40.64749  -73.97237
1   2595             Manhattan     Midtown  40.75362  -73.98377
2   3647             Manhattan       Harlem  40.80902  -73.94190
3   3831             Brooklyn  Clinton Hill  40.68514  -73.95976
4   5022             Manhattan   East Harlem  40.79851  -73.94399

      room_type  price  minimum_nights  number_of_reviews \
0   Private room   149                1                 9
1  Entire home/apt   225                1                45
2   Private room   150                3                 0
3  Entire home/apt    89                1               270
4  Entire home/apt    80               10                 9

      reviews_per_month  calculated_host_listings_count  availability_365
0                0.21                        6              365
1                0.38                        2              355
2                 NaN                        1              365
3                4.64                        1              194
4                0.10                        1               0
```

```
[36]: airbnb.describe()
```


```
[36]:      id  latitude  longitude  price  minimum_nights \
count  4.889500e+04  48895.000000  48895.000000  48895.000000  48895.000000
mean    1.901714e+07    40.728949   -73.952170    152.720687     7.029962
std     1.098311e+07     0.054530     0.046157    240.154170    20.510550
min     2.539000e+03    40.499790   -74.244420     0.000000     1.000000
```

25%	9.471945e+06	40.690100	-73.983070	69.000000	1.000000
50%	1.967728e+07	40.723070	-73.955680	106.000000	3.000000
75%	2.915218e+07	40.763115	-73.936275	175.000000	5.000000
max	3.648724e+07	40.913060	-73.712990	10000.000000	1250.000000


	number_of_reviews	reviews_per_month	calculated_host_listings_count	\
count	48895.000000	38843.000000		48895.000000
mean	23.274466	1.373221		7.143982
std	44.550582	1.680442		32.952519
min	0.000000	0.010000		1.000000
25%	1.000000	0.190000		1.000000
50%	5.000000	0.720000		1.000000
75%	24.000000	2.020000		2.000000
max	629.000000	58.500000		327.000000

	availability_365
count	48895.000000
mean	112.781327
std	131.622289
min	0.000000
25%	0.000000
50%	45.000000
75%	227.000000
max	365.000000


```
[37]: airbnb["price"].hist()
plt.show()
```



```
[38]: airbnb["calculated_host_listings_count"].hist()  
plt.show()
```



```
[39]: airbnb["availability_365"].hist()  
plt.show()
```





2.0.2 [5 pts] Plot total number_of_reviews per neighbourhood_group

```
[64]: answer = airbnb_avail.groupby("neighbourhood_group",
    ↪as_index=False)["number_of_reviews"].sum()
answer
```


```
[64]:  neighbourhood_group  number_of_reviews
0          Bronx          14080
1      Brooklyn      214091
2      Manhattan      186786
3          Queens       74420
4    Staten Island        7964
```

```
[66]: answer_graph = answer.plot.bar(x='neighbourhood_group', y='number_of_reviews',
    ↪rot=0)
```



2.0.3 [5 pts] Plot map of airbnbs throughout New York (if it gets too crowded take a subset of the data, and try to make it look nice if you can :)).

```
[40]: images_path = os.path.join('./', "images")
os.makedirs(images_path, exist_ok=True)
filename="newyorkcity.jpeg"
#import matplotlib.image as mpimg
newyork_img = mpimg.imread(os.path.join(images_path, filename))
ax = airbnb.plot(kind="scatter",x="longitude",y="latitude", alpha=0.4)
plt.imshow(newyork_img, extent=[-74.244420,-73.712990,40.499790,40.
↪913060],alpha=0.5)
plt.ylabel("Latitude",fontsize=14)
plt.xlabel("Longitude",fontsize=14)
plt.show()
```




2.0.4 [5 pts] Plot average price of room types who have availability greater than 180 days.

```
[41]: is_avail = airbnb["availability_365"] > 180
airbnb_avail = airbnb[is_avail]
airbnb_avail.head()
result = airbnb_avail.groupby("room_type", as_index=False)["price"].mean()
result
```

```
[41]:      room_type      price
0  Entire home/apt  248.870817
1    Private room   100.028192
2    Shared room    56.941909
```

```
[67]: result_graph = result.plot.bar(x='room_type', y='price', rot=0)
```




2.0.5 [5 pts] Plot correlation matrix

- which features have positive correlation?
- which features have negative correlation?

```
[63]: airbnb_corr_matrix = airbnb.corr()
      #from pandas.plotting import scatter_matrix
      attributes=["price", "number_of_reviews", "calculated_host_listings_count", "availability_365"]
      scatter_matrix(airbnb_corr_matrix[attributes], figsize=(15,10))
```


```
[63]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x1a2b7b5c50>,
              <matplotlib.axes._subplots.AxesSubplot object at 0x1a2bb14198>,
              <matplotlib.axes._subplots.AxesSubplot object at 0x1a2bb45748>,
              <matplotlib.axes._subplots.AxesSubplot object at 0x1a2bb75cf8>],
             [<matplotlib.axes._subplots.AxesSubplot object at 0x1a2bbb52e8>,
              <matplotlib.axes._subplots.AxesSubplot object at 0x1a2bbe3898>,
              <matplotlib.axes._subplots.AxesSubplot object at 0x1a2bc16e48>,
              <matplotlib.axes._subplots.AxesSubplot object at 0x1a2bc54470>],
             [<matplotlib.axes._subplots.AxesSubplot object at 0x1a2bc544a8>,
              <matplotlib.axes._subplots.AxesSubplot object at 0x1a2bcbbf98>,
              <matplotlib.axes._subplots.AxesSubplot object at 0x1a2bcf4588>,
              <matplotlib.axes._subplots.AxesSubplot object at 0x1a2bd27b38>],
             [<matplotlib.axes._subplots.AxesSubplot object at 0x1a2bd67128>,
              <matplotlib.axes._subplots.AxesSubplot object at 0x1a2bd966d8>],
```

```
<matplotlib.axes._subplots.AxesSubplot object at 0x1a2bdc8c88>,
<matplotlib.axes._subplots.AxesSubplot object at 0x1a2be06278>]],
dtype=object)
```



```
[44]: airbnb.plot(kind="scatter", x="reviews_per_month", y="number_of_reviews",
        alpha=0.1)
```


```
[44]: <matplotlib.axes._subplots.AxesSubplot at 0x1a22455278>
```



reviews_per_month and number_of_reviews: positive correlation

```
[62]: airbnb.plot(kind="scatter", x="price", y="minimum_nights",  
        alpha=0.1)
```


```
[62]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2dcc8a58>
```



price and minimum_nights: positive correlation

```
[48]: airbnb.plot(kind="scatter", x="price", y="availability_365",  
        alpha=0.1)
```


```
[48]: <matplotlib.axes._subplots.AxesSubplot at 0x1a259b22b0>
```



price and availability_365: no correlation

```
[49]: airbnb.plot(kind="scatter", x="calculated_host_listings_count", y="price",  
        alpha=0.1)
```


```
[49]: <matplotlib.axes._subplots.AxesSubplot at 0x1a26077390>
```

calculated_host_listings_count and price: positive correlation

```
[50]: airbnb.plot(kind="scatter", x="calculated_host_listings_count",  
      ↪y="number_of_reviews",  
      alpha=0.1)
```


```
[50]: <matplotlib.axes._subplots.AxesSubplot at 0x1a26292f60>
```



calculated_host_listings_count and number_of_reviews: positive correlation

```
[51]: airbnb.plot(kind="scatter", x="minimum_nights", y="availability_365",  
        alpha=0.1)
```



```
[51]: <matplotlib.axes._subplots.AxesSubplot at 0x1a26c5ddd8>
```



minimum_nights and availability_365: unclear correlation

```
[52]: airbnb.plot(kind="scatter", x="availability_365", y="number_of_reviews",
        alpha=0.1)
airbnb.plot(kind="scatter", x="availability_365", y="reviews_per_month",
        alpha=0.1)
```

```
[52]: <matplotlib.axes._subplots.AxesSubplot at 0x1a29f380b8>
```



availability_365 and number_of_reviews: unclear correlation

availability_365 and reviews_per_month: possibly negative correlation

Reviews per month and the number of reviews appears to have a positive correlation, which makes sense given that a higher number of reviews per month would end up increasing the total number of reviews. Other pairs of features such as minimum nights and price, host listings count and price, and host listings count and number of reviews look to only have slight positive correlation. It can be hard to tell because oftentimes the data will include a density of points that look like a straight line.

There don't appear to be any features that have an overt negative correlation. A case could be made for availability and the number of reviews per month, where a higher number of available days theoretically should correlate to a low number of reviews per month. After all, a higher number of available days means people aren't reserving the listing as much, so a lower number of guests should result in a low number of reviews per month. Given there are other factors that impact the number of reviews per month, like customer dissatisfaction or laziness, the correlation between availability and number of reviews per month is very hard to see on the graph.

3 [25 pts] Prepare the Data

3.0.1 [5 pts] Augment the dataframe with two other features which you think would be useful

```
[53]: airbnb["reviews_squared_per_month"] = airbnb["number_of_reviews"] *  
      ↪airbnb["reviews_per_month"]  
airbnb["nights_available_one_stay"] = airbnb["availability_365"] -  
      ↪airbnb["minimum_nights"]
```

3.0.2 [5 pts] Set aside 20% of the data as test test (80% train, 20% test).

```
[54]: from sklearn.model_selection import train_test_split  
airbnb = airbnb.drop(columns=['id', 'neighbourhood', 'latitude', 'longitude'])  
airbnb_x = airbnb.drop(columns=["price"])  
airbnb_y = airbnb["price"]  
x_train, x_test, y_train, y_test = train_test_split(airbnb_x, airbnb_y,  
      ↪test_size=0.2, random_state=42)
```

3.0.3 [5 pts] Impute any missing feature with a method of your choice, and briefly discuss why you chose this imputation method

```
[55]: incomplete_rows = airbnb[airbnb.isnull().any(axis=1)].head()  
incomplete_rows
```

```
[55]:   neighbourhood_group    room_type  price  minimum_nights  \  
2         Manhattan  Private room    150                3  
19        Manhattan  Entire home/apt    190                7
```

26	Manhattan	Private room	80	4
36	Brooklyn	Private room	35	60
38	Brooklyn	Private room	150	1

	number_of_reviews	reviews_per_month	calculated_host_listings_count	\
2	0	NaN		1
19	0	NaN		2
26	0	NaN		1
36	0	NaN		1
38	0	NaN		1

	availability_365	reviews_squared_per_month	nights_available_one_stay
2	365	NaN	362
19	249	NaN	242
26	0	NaN	-4
36	365	NaN	305
38	365	NaN	364

```
[56]: incomplete_rows["reviews_per_month"].fillna(0, inplace=True)
incomplete_rows.head()
```

```
[56]:
```

	neighbourhood_group	room_type	price	minimum_nights	\
2	Manhattan	Private room	150	3	
19	Manhattan	Entire home/apt	190	7	
26	Manhattan	Private room	80	4	
36	Brooklyn	Private room	35	60	
38	Brooklyn	Private room	150	1	

	number_of_reviews	reviews_per_month	calculated_host_listings_count	\
2	0	0.0		1
19	0	0.0		2
26	0	0.0		1
36	0	0.0		1
38	0	0.0		1

	availability_365	reviews_squared_per_month	nights_available_one_stay
2	365	NaN	362
19	249	NaN	242
26	0	NaN	-4
36	365	NaN	305
38	365	NaN	364

I chose to replace the null values with 0. I noticed that the null values occurred in the 'reviews_per_month' column and inferred that it was because the value in the 'number_of_reviews' column was 0 and whatever calculation outputted the number of reviews per month would output a null value. When the number of reviews was 0, I thought it was reasonable to set the number of reviews per month to 0 as well instead of a null value.

3.0.4 [10 pts] Code complete data pipeline using sklearn mixins

```
[57]: airbnb_num = x_train.drop(["neighbourhood_group", "room_type"], axis=1)
airbnb_num.head()
nights_ix, num_reviews_ix, reviewspm_ix, avail_ix = 1, 2, 3, 5

[58]: airbnb_imputer = SimpleImputer(strategy="constant", fill_value=0)

class Augment_Features1(BaseEstimator, TransformerMixin):
    def __init__(self, add_reviews_squared_per_month = True):
        self.add_reviews_squared_per_month = add_reviews_squared_per_month
    def fit(self, X, y=None):
        return self
    def transform(self, X):
        nights_available_one_stay = X[:, avail_ix] - X[:, nights_ix]
        if self.add_reviews_squared_per_month:
            reviews_squared_per_month = X[:, num_reviews_ix] * X[:,
↪reviewspm_ix]
            return np.c_[X, nights_available_one_stay,
↪reviews_squared_per_month]
        else:
            return np.c_[X, nights_available_one_stay]
#airbnb_attr_adder = Augment_Features1()
#airbnb_extra_attribs = airbnb_attr_adder.transform(x_train.values)

airbnb_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="constant", fill_value=1)),
    ('attribs_adder', Augment_Features1()),
    ('std_scaler', StandardScaler()),
])
airbnb_num_tr = airbnb_pipeline.fit_transform(airbnb_num)
num_features = list(airbnb_num)
cat_features = ["neighbourhood_group", "room_type"]

a_full_pipeline = ColumnTransformer([
    ("num", airbnb_pipeline, num_features),
    ("cat", OneHotEncoder(), cat_features)
])
airbnb_prepared = a_full_pipeline.fit_transform(x_train)
```

4 [15 pts] Fit a model of your choice

The task is to predict the price, you could refer to the housing example on how to train and evaluate your model using MSE. Provide both test and train set MSE values.

```
[59]: a_lin_reg = LinearRegression()
a_lin_reg.fit(airbnb_prepared, y_train)

a_data = x_test.iloc[:5]
a_labels = y_train.iloc[:5]
a_data_prepared = a_full_pipeline.transform(a_data)

print("Predictions:", a_lin_reg.predict(a_data_prepared))
print("Actual labels:", list(a_labels))
```

```
Predictions: [182.95703125  56.34765625 112.37890625 266.76953125 223.06640625]
Actual labels: [295, 70, 58, 75, 38]
```

```
[60]: a_preds = a_lin_reg.predict(airbnb_prepared)
a_mse = mean_squared_error(y_train, a_preds)
a_rmse = np.sqrt(a_mse)
a_rmse
```

```
[60]: 235.7809852797133
```

```
[ ]:
```