

# CS 188 Project 1

January 21, 2020

## 1 TODO: Applying the end-end ML steps to a different dataset.

We will apply what we've learnt to another dataset (airbnb dataset). We will predict airbnb price based on other features.

## 2 [25 pts] Visualizing Data

### 2.0.1 [5 pts] Load the data + statistics

- load the dataset
- display the first few rows of the data
- drop the following columns: name, host\_id, host\_name, last\_review
- display a summary of the statistics of the loaded data
- plot histograms for 3 features of your choice

```
[84]: import sys
assert sys.version_info >= (3, 5) # python>=3.5

import sklearn
assert sklearn.__version__ >= "0.20" # sklearn >= 0.20

import numpy as np #numerical package in python
import os
%matplotlib inline
import matplotlib.pyplot as plt #plotting package

# to make this notebook's output identical at every run
np.random.seed(42)

#matplotlib magic for inline figures
%matplotlib inline
import matplotlib # plotting library
import matplotlib.pyplot as plt

import os
import tarfile
```

```

import urllib

import pandas as pd

import matplotlib.image as mpimg

from pandas.plotting import scatter_matrix

from sklearn.model_selection import StratifiedShuffleSplit

from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer

from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder

from sklearn.base import BaseEstimator, TransformerMixin

from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_squared_error

from sklearn.model_selection import train_test_split

```

```

[85]: DATASET_PATH = os.path.join("datasets", "airbnb")

def load_airbnb_data(airbnb_path):
    csv_path = os.path.join(airbnb_path, "AB_NYC_2019.csv")
    return pd.read_csv(csv_path)
airbnb = load_airbnb_data(DATASET_PATH)
airbnb = airbnb.drop(columns=['name', 'host_id', 'host_name', 'last_review'])
airbnb.head()

```

```

[85]:      id  neighbourhood_group  neighbourhood  latitude  longitude  \
0   2539             Brooklyn    Kensington  40.64749   -73.97237
1   2595             Manhattan      Midtown  40.75362   -73.98377
2   3647             Manhattan      Harlem  40.80902   -73.94190
3   3831             Brooklyn  Clinton Hill  40.68514   -73.95976
4   5022             Manhattan    East Harlem  40.79851   -73.94399

      room_type  price  minimum_nights  number_of_reviews  \
0   Private room   149                1                 9
1  Entire home/apt   225                1                45
2   Private room   150                3                 0
3  Entire home/apt    89                1               270
4  Entire home/apt    80               10                 9

```

	reviews_per_month	calculated_host_listings_count	availability_365
0	0.21	6	365
1	0.38	2	355
2	NaN	1	365
3	4.64	1	194
4	0.10	1	0

```
[86]: airbnb.describe()
```

```
[86]:
```

	id	latitude	longitude	price	minimum_nights \
count	4.889500e+04	48895.000000	48895.000000	48895.000000	48895.000000
mean	1.901714e+07	40.728949	-73.952170	152.720687	7.029962
std	1.098311e+07	0.054530	0.046157	240.154170	20.510550
min	2.539000e+03	40.499790	-74.244420	0.000000	1.000000
25%	9.471945e+06	40.690100	-73.983070	69.000000	1.000000
50%	1.967728e+07	40.723070	-73.955680	106.000000	3.000000
75%	2.915218e+07	40.763115	-73.936275	175.000000	5.000000
max	3.648724e+07	40.913060	-73.712990	10000.000000	1250.000000


  

	number_of_reviews	reviews_per_month	calculated_host_listings_count \
count	48895.000000	38843.000000	48895.000000
mean	23.274466	1.373221	7.143982
std	44.550582	1.680442	32.952519
min	0.000000	0.010000	1.000000
25%	1.000000	0.190000	1.000000
50%	5.000000	0.720000	1.000000
75%	24.000000	2.020000	2.000000
max	629.000000	58.500000	327.000000


  

	availability_365
count	48895.000000
mean	112.781327
std	131.622289
min	0.000000
25%	0.000000
50%	45.000000
75%	227.000000
max	365.000000


```
[87]: airbnb["price"].hist()
plt.show()
```



```
[88]: airbnb["calculated_host_listings_count"].hist()  
plt.show()
```



```
[89]: airbnb["availability_365"].hist()  
plt.show()
```




## 2.0.2 [5 pts] Plot total number\_of\_reviews per neighbourhood\_group

```
[90]: answer = airbnb.groupby("neighbourhood_group",  
    ↳as_index=False)["number_of_reviews"].sum()  
answer
```


```
[90]:  neighbourhood_group  number_of_reviews  
0           Bronx         28371  
1       Brooklyn     486574  
2       Manhattan     454569  
3         Queens     156950  
4  Staten Island      11541
```

```
[91]: answer_graph = answer.plot.bar(x='neighbourhood_group', y='number_of_reviews',  
    ↳rot=0)
```



**2.0.3 [5 pts]** Plot map of airbnbs throughout New York (if it gets too crowded take a subset of the data, and try to make it look nice if you can :)).

```
[92]: images_path = os.path.join('./', "images")
os.makedirs(images_path, exist_ok=True)
filename="newyorkcity.jpeg"
#import matplotlib.image as mpimg
newyork_img = mpimg.imread(os.path.join(images_path, filename))
ax = airbnb.plot(kind="scatter",x="longitude",y="latitude", alpha=0.4)
plt.imshow(newyork_img, extent=[-74.244420,-73.712990,40.499790,40.
↪913060],alpha=0.5)
plt.ylabel("Latitude",fontsize=14)
plt.xlabel("Longitude",fontsize=14)
plt.show()
```




2.0.4 [5 pts] Plot average price of room types who have availability greater than 180 days.

```
[93]: is_avail = airbnb["availability_365"] > 180
airbnb_avail = airbnb[is_avail]
airbnb_avail.head()
result = airbnb_avail.groupby("room_type", as_index=False)["price"].mean()
result
```

```
[93]:      room_type      price
0  Entire home/apt  248.870817
1    Private room  100.028192
2    Shared room   56.941909
```

```
[94]: result_graph = result.plot.bar(x='room_type', y='price', rot=0)
```



### 2.0.5 [5 pts] Plot correlation matrix

- which features have positive correlation?
- which features have negative correlation?

```
[95]: airbnb_corr_matrix = airbnb.corr()
      #from pandas.plotting import scatter_matrix
      attributes=["price", "number_of_reviews", "calculated_host_listings_count", "availability_365"]
      scatter_matrix(airbnb_corr_matrix[attributes], figsize=(15,10))
```


```
[95]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x1a28bed0b8>,
               <matplotlib.axes._subplots.AxesSubplot object at 0x1a274106d8>,
               <matplotlib.axes._subplots.AxesSubplot object at 0x1a27679c88>,
               <matplotlib.axes._subplots.AxesSubplot object at 0x1a2990a278>],
             [<matplotlib.axes._subplots.AxesSubplot object at 0x1a29939828>,
               <matplotlib.axes._subplots.AxesSubplot object at 0x1a2a12cdd8>,
               <matplotlib.axes._subplots.AxesSubplot object at 0x1a303ff5f8>,
               <matplotlib.axes._subplots.AxesSubplot object at 0x1a303cc5f8>],
             [<matplotlib.axes._subplots.AxesSubplot object at 0x1a303cc780>,
               <matplotlib.axes._subplots.AxesSubplot object at 0x1a2df48588>,
               <matplotlib.axes._subplots.AxesSubplot object at 0x1a2dfceba8>,
               <matplotlib.axes._subplots.AxesSubplot object at 0x1a2747fd30>],
             [<matplotlib.axes._subplots.AxesSubplot object at 0x1a26acf5c0>,
               <matplotlib.axes._subplots.AxesSubplot object at 0x1a26a8d898>],
             ]
```



```

<matplotlib.axes._subplots.AxesSubplot object at 0x1a26e8e6d8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x1a2520feb8>]],
dtype=object)

```



```


[96]: airbnb.plot(kind="scatter", x="reviews_per_month", y="number_of_reviews",
        alpha=0.1)

```

```

[96]: <matplotlib.axes._subplots.AxesSubplot at 0x1a27251748>


```



reviews\_per\_month and number\_of\_reviews: positive correlation

```
[97]: airbnb.plot(kind="scatter", x="price", y="minimum_nights",  
        alpha=0.1)
```


```
[97]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2e3db438>
```



price and minimum\_nights: positive correlation

```
[98]: airbnb.plot(kind="scatter", x="price", y="availability_365",  
              alpha=0.1)
```


```
[98]: <matplotlib.axes._subplots.AxesSubplot at 0x1a28696780>
```



price and availability\_365: no correlation

```
[99]: airbnb.plot(kind="scatter", x="calculated_host_listings_count", y="price",  
         alpha=0.1)
```

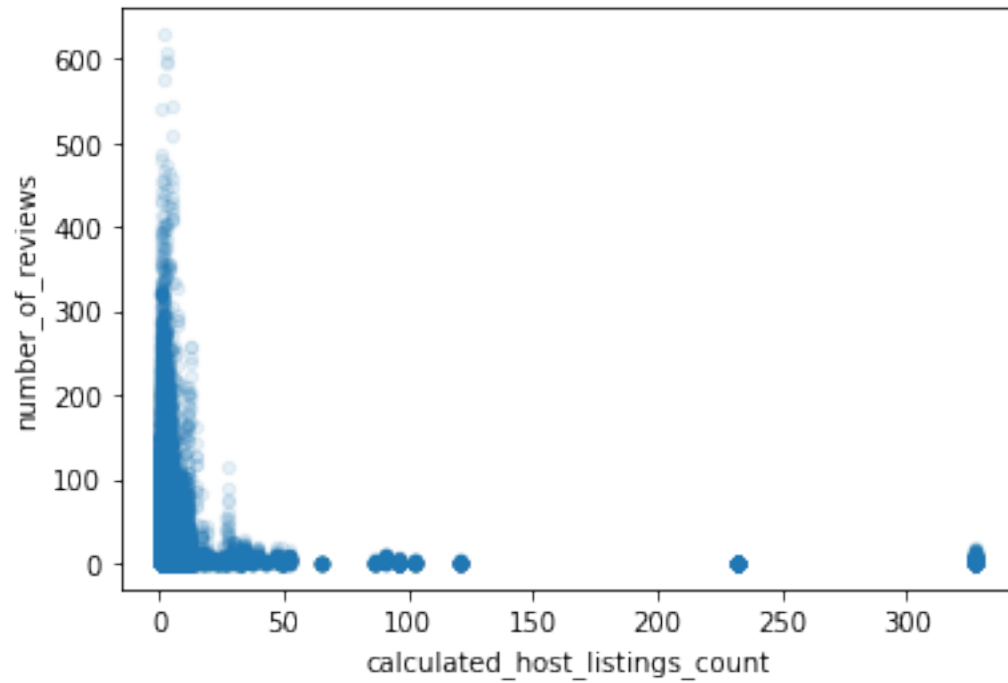
```
[99]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2e25a908>
```



calculated\_host\_listings\_count and price: positive correlation

```
[100]: airbnb.plot(kind="scatter", x="calculated_host_listings_count",  
        ↪y="number_of_reviews",  
        alpha=0.1)
```


```
[100]: <matplotlib.axes._subplots.AxesSubplot at 0x1a289a7d68>
```



calculated\_host\_listings\_count and number\_of\_reviews: positive correlation

```
[101]: airbnb.plot(kind="scatter", x="minimum_nights", y="availability_365",  
          alpha=0.1)
```



```
[101]: <matplotlib.axes._subplots.AxesSubplot at 0x1a2f94e588>
```



minimum\_nights and availability\_365: unclear correlation

```
[102]: airbnb.plot(kind="scatter", x="availability_365", y="number_of_reviews",
        alpha=0.1)
airbnb.plot(kind="scatter", x="availability_365", y="reviews_per_month",
        alpha=0.1)
```

```
[102]: <matplotlib.axes._subplots.AxesSubplot at 0x1a30e7c080>
```



availability\_365 and number\_of\_reviews: unclear correlation

availability\_365 and reviews\_per\_month: possibly negative correlation



Reviews per month and the number of reviews appears to have a positive correlation, which makes sense given that a higher number of reviews per month would end up increasing the total number of reviews. Other pairs of features such as minimum nights and price, host listings count and price, and host listings count and number of reviews look to only have slight positive correlation. It can be hard to tell because oftentimes the data will include a density of points that look like a straight line.

There don't appear to be any features that have an overt negative correlation. A case could be made for availability and the number of reviews per month, where a higher number of available days theoretically should correlate to a low number of reviews per month. After all, a higher number of available days means people aren't reserving the listing as much, so a lower number of guests should result in a low number of reviews per month. Given there are other factors that impact the number of reviews per month, like customer dissatisfaction or laziness, the correlation between availability and number of reviews per month is very hard to see on the graph.

### 3 [25 pts] Prepare the Data

3.0.1 [5 pts] Augment the dataframe with two other features which you think would be useful

```
[103]: airbnb["reviews_squared_per_month"] = airbnb["number_of_reviews"] *  
        ↪airbnb["reviews_per_month"]  
airbnb["nights_available_one_stay"] = airbnb["availability_365"] -  
        ↪airbnb["minimum_nights"]
```

3.0.2 [5 pts] Set aside 20% of the data as test test (80% train, 20% test).

```
[104]: airbnb = airbnb.drop(columns=['id', 'neighbourhood', 'latitude', 'longitude'])  
airbnb_x = airbnb.drop(columns=["price"])  
airbnb_y = airbnb["price"]  
x_train, x_test, y_train, y_test = train_test_split(airbnb_x, airbnb_y,  
        ↪test_size=0.2, random_state=42)
```

3.0.3 [5 pts] Impute any missing feature with a method of your choice, and briefly discuss why you chose this imputation method

```
[105]: incomplete_rows = airbnb[airbnb.isnull().any(axis=1)].head()  
incomplete_rows
```

```
[105]:
```

	neighbourhood_group	room_type	price	minimum_nights	\
2	Manhattan	Private room	150	3	
19	Manhattan	Entire home/apt	190	7	
26	Manhattan	Private room	80	4	

36	Brooklyn	Private room	35	60
38	Brooklyn	Private room	150	1

	number_of_reviews	reviews_per_month	calculated_host_listings_count	\
2	0	NaN		1
19	0	NaN		2
26	0	NaN		1
36	0	NaN		1
38	0	NaN		1

	availability_365	reviews_squared_per_month	nights_available_one_stay
2	365	NaN	362
19	249	NaN	242
26	0	NaN	-4
36	365	NaN	305
38	365	NaN	364

```
[106]: incomplete_rows["reviews_per_month"].fillna(0, inplace=True)
incomplete_rows.head()
```

```
[106]:
```

	neighbourhood_group	room_type	price	minimum_nights	\
2	Manhattan	Private room	150	3	
19	Manhattan	Entire home/apt	190	7	
26	Manhattan	Private room	80	4	
36	Brooklyn	Private room	35	60	
38	Brooklyn	Private room	150	1	

	number_of_reviews	reviews_per_month	calculated_host_listings_count	\
2	0	0.0		1
19	0	0.0		2
26	0	0.0		1
36	0	0.0		1
38	0	0.0		1

	availability_365	reviews_squared_per_month	nights_available_one_stay
2	365	NaN	362
19	249	NaN	242
26	0	NaN	-4
36	365	NaN	305
38	365	NaN	364

I chose to replace the null values with 0. I noticed that the null values occurred in the 'reviews\_per\_month' column and inferred that it was because the value in the 'number\_of\_reviews' column was 0 and whatever calculation outputted the number of reviews per month would output a null value. When the number of reviews was 0, I thought it was reasonable to set the number of reviews per month to 0 as well instead of a null value.

### 3.0.4 [10 pts] Code complete data pipeline using sklearn mixins

```
[107]: airbnb_num = x_train.drop(["neighbourhood_group", "room_type"], axis=1)
airbnb_num.head()
nights_ix, num_reviews_ix, reviewspm_ix, avail_ix = 1, 2, 3, 5

[108]: airbnb_imputer = SimpleImputer(strategy="constant", fill_value=0)

class Augment_Features1(BaseEstimator, TransformerMixin):
    def __init__(self, add_reviews_squared_per_month = True):
        self.add_reviews_squared_per_month = add_reviews_squared_per_month
    def fit(self, X, y=None):
        return self
    def transform(self, X):
        nights_available_one_stay = X[:, avail_ix] - X[:, nights_ix]
        if self.add_reviews_squared_per_month:
            reviews_squared_per_month = X[:, num_reviews_ix] * X[:,
↪reviewspm_ix]
            return np.c_[X, nights_available_one_stay,
↪reviews_squared_per_month]
        else:
            return np.c_[X, nights_available_one_stay]
#airbnb_attr_adder = Augment_Features1()
#airbnb_extra_attribs = airbnb_attr_adder.transform(x_train.values)

airbnb_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="constant", fill_value=1)),
    ('attribs_adder', Augment_Features1()),
    ('std_scaler', StandardScaler()),
])
airbnb_num_tr = airbnb_pipeline.fit_transform(airbnb_num)
num_features = list(airbnb_num)
cat_features = ["neighbourhood_group", "room_type"]

a_full_pipeline = ColumnTransformer([
    ("num", airbnb_pipeline, num_features),
    ("cat", OneHotEncoder(), cat_features)
])
airbnb_prepared = a_full_pipeline.fit_transform(x_train)
```

## 4 [15 pts] Fit a model of your choice

The task is to predict the price, you could refer to the housing example on how to train and evaluate your model using MSE. Provide both test and train set MSE values.

```
[109]: a_lin_reg = LinearRegression()
a_lin_reg.fit(airbnb_prepared, y_train)

a_data = x_test.iloc[:5]
a_labels = y_train.iloc[:5]
a_data_prepared = a_full_pipeline.transform(a_data)

print("Predictions:", a_lin_reg.predict(a_data_prepared))
print("Actual labels:", list(a_labels))
```

```
Predictions: [182.95703125  56.34765625 112.37890625 266.76953125 223.06640625]
Actual labels: [295, 70, 58, 75, 38]
```

```
[111]: a_preds = a_lin_reg.predict(airbnb_prepared)
a_mse = mean_squared_error(y_train, a_preds)
a_rmse = np.sqrt(a_mse)
print("Mean squared loss:", a_rmse)
```

```
Mean squared loss: 235.7809852797133
```

```
[ ]:
```